

Capstone: Instant Health Alert System – Final Submission

EMR instances

- EMR cluster (with Spark, Hive, Sqoop)

Cluster: SparkwithHiveSqoopcluster Starting

SummaryApplication user interfacesMonitoringHardwareConfigurationsEventsStepsBootstrap actions

Summary

ID: j-1RGRZKURC1L6Q

Creation date: 2023-03-28 14:00 (UTC+5:30)

Elapsed time: 0 seconds

After last step completes: Cluster waits

Termination protection: Off [Change](#)

Tags: -- [View All / Edit](#)

Master public DNS: --

Application user interfaces

Persistent user interfaces [🔗](#): --

On-cluster user -- interfaces [🔗](#):

Security and access

Key name: MyNew_KeyValue

EC2 instance profile: EMR_EC2_DefaultRole

EMR role: EMR_DefaultRole

Configuration details

Release label: emr-5.30.1

Hadoop distribution: Amazon 2.8.5

Applications: Spark 2.4.5, JupyterHub 1.1.0, Zeppelin 0.8.2, Livy 0.7.0, Hive 2.3.6, HCatalog 2.3.6, Sqoop 1.4.7

Log URI: s3://aws-logs-545120555452-us-east-1/elasticmapreduce/ [📁](#)

EMRFS consistent view: Disabled

Custom AMI ID: --

Network and hardware

Availability zone: --

Subnet ID: [subnet-007e940c0eb58c44b](#) [🔗](#)

Master: Provisioning 1 m4.xlarge

Core: --

Task: --

Cluster scaling: Not enabled

Auto-termination: Not enabled

EMR Hardware Configuration (with 1 master node and 2 core nodes)

Cluster: SparkwithHiveSqoopcluster Terminated Terminated by user request

SummaryApplication user interfacesMonitoringHardwareConfigurationsEventsStepsBootstrap actions

Add task instance group

Instance groups

Filter: 2 instance groups (all loaded) [🔄](#)

ID	Status	Node type & name	Instance type	Instance count	Purchasing option
ig-14KVR95NI797O	Terminated (1 Requested)	MASTER Master - 1	m4.xlarge 4 vCore, 16 GiB memory, EBS only storage EBS Storage: 70 GiB	0 Instances	On-demand ⓘ
ig-H78HMFXPML5O	Terminated	CORE Core - 2	m4.xlarge 4 vCore, 16 GiB memory, EBS only storage EBS Storage: 64 GiB	0 Instances	On-demand ⓘ

- HBase Cluster

Cluster: Hbasecluster Waiting Cluster ready to run steps.

Summary

Application user interfaces

Monitoring

Hardware

Configurations

Events

Steps

Bootstrap actions

Summary

Configuration details

Application user interfaces

Network and hardware

Security and access

Summary

Configuration details

Application user interfaces

Network and hardware

Security and access

Summary

Configuration details

Application user interfaces

Network and hardware

Security and access

Summary

Configuration details

Application user interfaces

Network and hardware

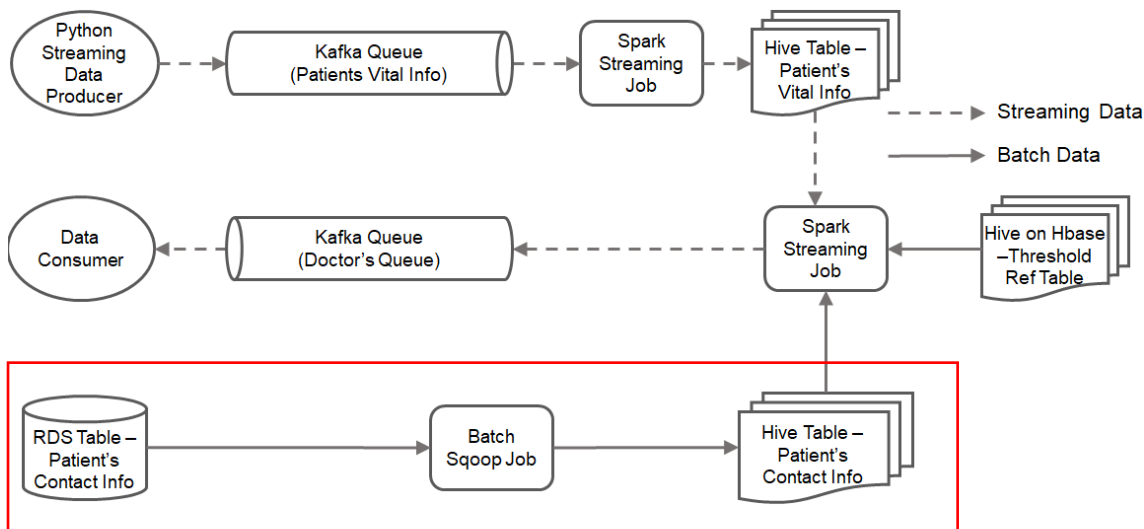
Security and access

EC2 instance

Covered in PART 3

PART 1:

We will be importing the patient contact info from the RDS using Sqoop and loading the data further into Hive table



1. Import Patient Contact Info records to HDFS using Sqoop
2. Creating an external Hive table for storing Patient's Contact Info

SQOOP SETUP

Following steps are followed to setup Sqoop on EMR Cluster

1. To install the MySQL connector jar file.

```
wget https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
```

```
[hadoop@ip-172-31-83-130 ~]$ wget https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
--2023-03-25 07:08:33-- https://de-mysql-connector.s3.amazonaws.com/mysql-connector-java-8.0.25.tar.gz
Resolving de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com)... 54.231.130.1, 3.5.16.186, 3.5.20.112, ...
Connecting to de-mysql-connector.s3.amazonaws.com (de-mysql-connector.s3.amazonaws.com) [54.231.130.1]:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 4079310 (3.9M) [application/x-gzip]
Saving to: 'mysql-connector-java-8.0.25.tar.gz'

100%[=====>] 4,079,310 20.6MB/s in 0.2s

2023-03-25 07:08:34 (20.6 MB/s) - 'mysql-connector-java-8.0.25.tar.gz' saved [4079310/4079310]
```

2. Extract the MySQL connector tar file

```
tar -xvf mysql-connector-java-8.0.25.tar.gz
```

```
[hadoop@ip-172-31-83-130 ~]$ tar -xvf mysql-connector-java-8.0.25.tar.gz
mysql-connector-java-8.0.25/
mysql-connector-java-8.0.25/src/
mysql-connector-java-8.0.25/src/build/
mysql-connector-java-8.0.25/src/build/java/
mysql-connector-java-8.0.25/src/build/java/documentation/
mysql-connector-java-8.0.25/src/build/java/instrumentation/
mysql-connector-java-8.0.25/src/build/misc/
mysql-connector-java-8.0.25/src/build/misc/debian.in/
mysql-connector-java-8.0.25/src/build/misc/debian.in/source/
mysql-connector-java-8.0.25/src/demo/
mysql-connector-java-8.0.25/src/demo/java/
mysql-connector-java-8.0.25/src/demo/java/demo/
mysql-connector-java-8.0.25/src/demo/java/demo/x/
mysql-connector-java-8.0.25/src/demo/java/demo/x/devapi/
mysql-connector-java-8.0.25/src/generated/
mysql-connector-java-8.0.25/src/generated/java/
mysql-connector-java-8.0.25/src/generated/java/com/
mysql-connector-java-8.0.25/src/generated/java/com/mysql/
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/x/
mysql-connector-java-8.0.25/src/generated/java/com/mysql/cj/x/protobuf/
mysql-connector-java-8.0.25/src/legacy/
mysql-connector-java-8.0.25/src/legacy/java/
mysql-connector-java-8.0.25/src/legacy/java/com/
mysql-connector-java-8.0.25/src/legacy/java/com/mysql/
mysql-connector-java-8.0.25/src/legacy/java/com/mysql/jdbc/
```

3. Go to the MySQL Connector directory created in the previous step and copy it to the Sqoop library to complete the installation.

```
cd mysql-connector-java-8.0.25/
sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/
```

```
[hadoop@ip-172-31-83-130 ~]$ cd mysql-connector-java-8.0.25/
[hadoop@ip-172-31-83-130 mysql-connector-java-8.0.25]$ sudo cp mysql-connector-java-8.0.25.jar /usr/lib/sqoop/lib/
```

4. Set up MySQL on your EMR cluster (Inside this folder mysql-connector-java-8.0.25)

```
mysql_secure_installation
```

Enter current password for root (enter for none): ENTER

Set root password [Y/n] Y
New password: 123
Re-enter password: 123
Remove anonymous users [Y/n] Y
Disallow root login remotely [Y/n] n
Remove test database and access to it [Y/n] Y
Reload privilege tables now [Y/n] Y

```
[hadoop@ip-172-31-83-130 mysql-connector-java-8.0.25]$ mysql_secure_installation
```

NOTE: RUNNING ALL PARTS OF THIS SCRIPT IS RECOMMENDED FOR ALL MariaDB
SERVERS IN PRODUCTION USE! PLEASE READ EACH STEP CAREFULLY!

In order to log into MariaDB to secure it, we'll need the current
password for the root user. If you've just installed MariaDB, and
you haven't set the root password yet, the password will be blank,
so you should just press enter here.

Enter current password for root (enter for none):
OK, successfully used password, moving on...

Setting the root password ensures that nobody can log into the MariaDB
root user without the proper authorisation.

Set root password? [Y/n] Y
New password:
Re-enter new password:
Password updated successfully!
Reloading privilege tables..
... Success!

By default, a MariaDB installation has an anonymous user, allowing anyone
to log into MariaDB without having to have a user account created for
them. This is intended only for testing, and to make the installation
go a bit smoother. You should remove them before moving into a
production environment.

Remove anonymous users? [Y/n] Y
... Success!

Normally, root should only be allowed to connect from 'localhost'. This
ensures that someone cannot guess at the root password from the network.

Disallow root login remotely? [Y/n] n
... skipping.

By default, MariaDB comes with a database named 'test' that anyone can
access. This is also intended only for testing, and should be removed

By default, MariaDB comes with a database named 'test' that anyone can
access. This is also intended only for testing, and should be removed
before moving into a production environment.

Remove test database and access to it? [Y/n] Y
- Dropping test database...
... Success!
- Removing privileges on test database...
... Success!

Reloading the privilege tables will ensure that all changes made so far
will take effect immediately.

Reload privilege tables now? [Y/n] Y
... Success!

Cleaning up...

All done! If you've completed all of the above steps, your MariaDB
installation should now be secure.

Thanks for using MariaDB!

6. With this, MySQL setup is done. Now, we can access the MySQL shell. Enter the following command, type 123 when the password prompt comes up, and finally, press Enter.

```
mysql -u root -p
```

```
[hadoop@ip-172-31-83-130 mysql-connector-java-8.0.25]$ mysql -u root -p
Enter password:
Welcome to the MariaDB monitor.  Commands end with ; or \g.
Your MariaDB connection id is 66
Server version: 5.5.68-MariaDB MariaDB Server

Copyright (c) 2000, 2018, Oracle, MariaDB Corporation Ab and others.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

MariaDB [(none)]> GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' identified by '123'
```

7. Inside MariaDB (MariaDB >)

Following queries need to be run for granting all privileges to the root user.

```
GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' identified by '123' WITH GRANT OPTION;
flush privileges;
exit;
```

```
MariaDB [(none)]> GRANT ALL PRIVILEGES ON *.* TO 'root'@'%' identified by '123'
WITH GRANT OPTION;
Query OK, 0 rows affected (0.00 sec)

MariaDB [(none)]> flush privileges;
Query OK, 0 rows affected (0.00 sec)

MariaDB [(none)]> exit;
Bye
```

8. Restart the MySQL service to finish setting up MySQL. (Inside this folder mysql-connector-java-8.0.25)

```
sudo service mariadb restart
```

```
[hadoop@ip-172-31-83-130 mysql-connector-java-8.0.25]$ sudo service mariadb rest
art
Redirecting to /bin/systemctl restart mariadb.service
```

9. Change the directory (come outside mysql-connector-java-8.0.25 folder)

```
cd ..
```

Sqoop Commands

1. Import data to HDFS

```
sqoop import --connect jdbc:mysql://upgradetest.cyaieic9bmnf.us-east-
1.rds.amazonaws.com/testdatabase --table patients_information --username student --password
STUDENT123 --target-dir /user/livy/patient_contact_info -m 1
```

```
[root@ip-172-31-83-130 ~]# sqoop import --connect jdbc:mysql://upgraddetest.cyai
elc9bmnf.us-east-1.rds.amazonaws.com/testdatabase --table patients_information -
-username student --password STUDENT123 --target-dir /user/livy/patient_contact_
info -m 1
```

```
Other local map tasks=1
Total time spent by all maps in occupied slots (ms)=161328
Total time spent by all reduces in occupied slots (ms)=0
Total time spent by all map tasks (ms)=3361
Total vcore-milliseconds taken by all map tasks=3361
Total megabyte-milliseconds taken by all map tasks=5162496
Map-Reduce Framework
Map input records=5
Map output records=5
Input split bytes=87
Spilled Records=0
Failed Shuffles=0
Merged Map outputs=0
GC time elapsed (ms)=67
CPU time spent (ms)=1890
Physical memory (bytes) snapshot=261730304
Virtual memory (bytes) snapshot=3281002496
Total committed heap usage (bytes)=247463936
File Input Format Counters
Bytes Read=0
File Output Format Counters
Bytes Written=230
23/03/25 07:11:39 INFO mapreduce.ImportJobBase: Transferred 230 bytes in 20.9571
seconds (10.9748 bytes/sec)
23/03/25 07:11:39 INFO mapreduce.ImportJobBase: Retrieved 5 records.
```

2. View the list of files in HDFS target directory

```
hadoop fs -ls /user/livy/patient_contact_info
```

```
[root@ip-172-31-83-130 ~]# hadoop fs -ls /user/livy/patient_contact_info
Found 2 items
-rw-r--r--  1 root livy          0 2023-03-25 07:11 /user/livy/patient_contact_
info/_SUCCESS
-rw-r--r--  1 root livy      230 2023-03-25 07:11 /user/livy/patient_contact_
info/part-m-00000
```

3. View the imported contents in HDFS file

```
hadoop fs -cat /user/livy/patient_contact_info/part-m-00000
```

```
[root@ip-172-31-83-130 ~]# hadoop fs -cat /user/livy/patient_contact_info/part-m
-00000
1,Alex S,XDC test Address,8982739282,1,23,null
2,Sammy A,New Building Address,2382739282,2,45,null
3,Karan C,Aws Address,8923739282,3,56,null
4,Dara M,India Address,2182739282,4,67,null
5,Pam,ABC test Address,4982739282,5,72,null
```

Hive table creation (for Patients_Contact_Info)

- Open Hive shell.

```
^C[hadoop@ip-172-31-82-68 ~]$ hive
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.
properties Async: true
```

- Create a database *patient_health_care*

```
create database if not exists patient_health_care;
```

```
hive> create database if not exists patient_health_care;  
OK  
Time taken: 0.855 seconds
```

- Use database patient_health_care

```
use patient_health_care;
```

```
hive> use patient_health_care;  
OK  
Time taken: 0.046 seconds
```

- Create external table named *Patients_Contact_Info*

```
CREATE EXTERNAL TABLE IF NOT EXISTS Patients_Contact_Info (  
  patientid int,  
  patientname string,  
  patientaddress string,  
  phone_number string,  
  admitted_ward int,  
  age int,  
  other_details string  
)  
row format delimited  
fields terminated by ','  
lines terminated by '\n'  
location '/user/livy/patient_contact_info';
```

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS Patients_Contact_Info (  
  >   patientid int,  
  >   patientname string,  
  >   patientaddress string,  
  >   phone_number string,  
  >   admitted_ward int,  
  >   age int,  
  >   other_details string  
  > )  
  > row format delimited  
  > fields terminated by ','  
  > lines terminated by '\n'  
  > location '/user/livy/patient_contact_info';  
OK  
Time taken: 0.318 seconds
```

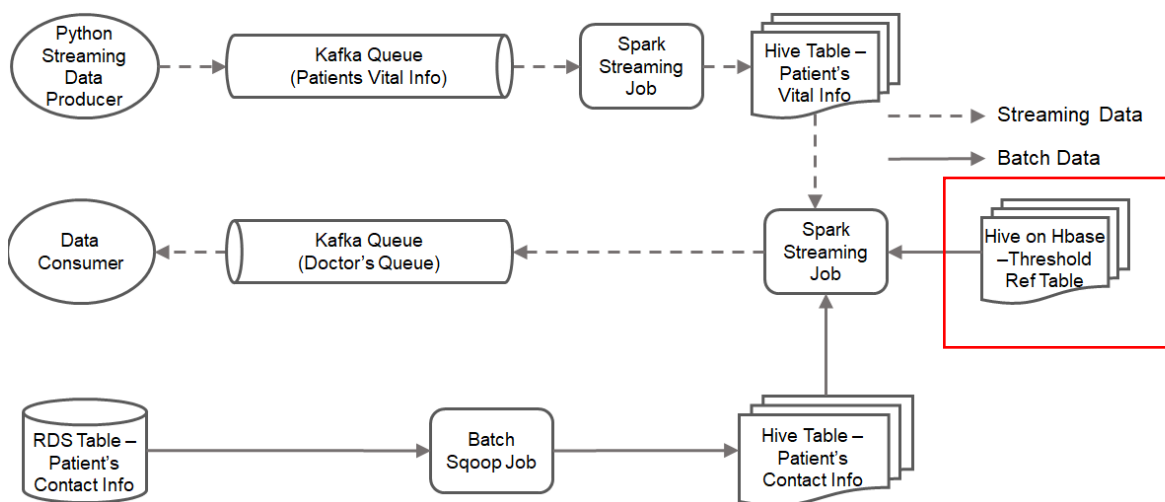
- View the records in *Patients_Contact_Info* table

```
select * from Patients_Contact_Info;
```

```
hive> select * from Patients_Contact_Info;
OK
patients_contact_info.patientid patients_contact_info.patientname patients
_contact_info.patientaddress patients_contact_info.phone_number patients
_contact_info.admitted_ward patients_contact_info.age patients_contact
_info.other_details
1 Alex S XDC test Address 8982739282 1 23 null
2 Sammy A New Building Address 2382739282 2 45 null
3 Karan C Aws Address 8923739282 3 56 null
4 Dara M India Address 2182739282 4 67 null
5 Pam ABC test Address 4982739282 5 72 null
Time taken: 1.541 seconds, Fetched: 5 row(s)
```

PART 2:

Create an HBase table to store threshold reference information and create a hive external table on top of this HBase table



1. Create a HBase table named **threshold_ref** with 3 column families: attribute, limit, alert
2. Insert 12 records in this HBase table
3. Set up Hive-HBase integration (since HBase and Hive are on separate clusters)
4. Create a Hive external table named **Threshold_Reference_Table** on top of HBase table

Navigate to HBase shell using below commands:

```
sudo -i
hbase shell
```

```
[root@ip-172-31-92-60 ~]# hbase shell
HBase Shell
Use "help" to get list of supported commands.
Use "exit" to quit this interactive shell.
Version 1.4.13, rUnknown, Fri Apr 17 15:18:24 UTC 2020
```

Create threshold_ref table in HBase

```
create 'threshold_ref','attribute','limit','alert';
```



```
hbase(main):001:0> create 'threshold_ref','attribute','limit','alert'  
0 row(s) in 1.6340 seconds
```

Insert 12 records into HBase table

```
put 'threshold_ref', '1', 'attribute:attribute', 'heartBeat';  
put 'threshold_ref', '1', 'limit:low_age_limit', '0';  
put 'threshold_ref', '1', 'limit:high_age_limit', '40';  
put 'threshold_ref', '1', 'limit:low_value', '0';  
put 'threshold_ref', '1', 'limit:high_value', '69';  
put 'threshold_ref', '1', 'alert:alert_flag', '1';  
put 'threshold_ref', '1', 'alert:alert_message', 'Low Heart Rate than Normal';  
  
put 'threshold_ref', '2', 'attribute:attribute', 'heartBeat';  
put 'threshold_ref', '2', 'limit:low_age_limit', '0';  
put 'threshold_ref', '2', 'limit:high_age_limit', '40';  
put 'threshold_ref', '2', 'limit:low_value', '70';  
put 'threshold_ref', '2', 'limit:high_value', '78';  
put 'threshold_ref', '2', 'alert:alert_flag', '0';  
put 'threshold_ref', '2', 'alert:alert_message', 'Normal';  
  
put 'threshold_ref', '3', 'attribute:attribute', 'heartBeat';  
put 'threshold_ref', '3', 'limit:low_age_limit', '0';  
put 'threshold_ref', '3', 'limit:high_age_limit', '40';  
put 'threshold_ref', '3', 'limit:low_value', '79';  
put 'threshold_ref', '3', 'limit:high_value', '9999';  
put 'threshold_ref', '3', 'alert:alert_flag', '1';  
put 'threshold_ref', '3', 'alert:alert_message', 'Higher Heart Rate than Normal';  
  
put 'threshold_ref', '4', 'attribute:attribute', 'bp';  
put 'threshold_ref', '4', 'limit:low_age_limit', '0';  
put 'threshold_ref', '4', 'limit:high_age_limit', '40';  
put 'threshold_ref', '4', 'limit:low_value', '0';  
put 'threshold_ref', '4', 'limit:high_value', '160';  
put 'threshold_ref', '4', 'alert:alert_flag', '1';  
put 'threshold_ref', '4', 'alert:alert_message', 'Low BP than Normal';  
  
put 'threshold_ref', '5', 'attribute:attribute', 'bp';  
put 'threshold_ref', '5', 'limit:low_age_limit', '0';  
put 'threshold_ref', '5', 'limit:high_age_limit', '40';  
put 'threshold_ref', '5', 'limit:low_value', '161';  
put 'threshold_ref', '5', 'limit:high_value', '220';  
put 'threshold_ref', '5', 'alert:alert_flag', '0';  
put 'threshold_ref', '5', 'alert:alert_message', 'Normal';  
  
put 'threshold_ref', '6', 'attribute:attribute', 'bp';  
put 'threshold_ref', '6', 'limit:low_age_limit', '0';  
put 'threshold_ref', '6', 'limit:high_age_limit', '40';  
put 'threshold_ref', '6', 'limit:low_value', '221';  
put 'threshold_ref', '6', 'limit:high_value', '9999';  
put 'threshold_ref', '6', 'alert:alert_flag', '1';  
put 'threshold_ref', '6', 'alert:alert_message', 'Higer BP than Normal';  
  
put 'threshold_ref', '7', 'attribute:attribute', 'heartBeat';  
put 'threshold_ref', '7', 'limit:low_age_limit', '41';  
put 'threshold_ref', '7', 'limit:high_age_limit', '100';
```

```
put 'threshold_ref', '7', 'limit:low_value', '0';
put 'threshold_ref', '7', 'limit:high_value', '65';
put 'threshold_ref', '7', 'alert:alert_flag', '1';
put 'threshold_ref', '7', 'alert:alert_message', 'Low Heart Rate than Normal';

put 'threshold_ref', '8', 'attribute:attribute', 'heartBeat';
put 'threshold_ref', '8', 'limit:low_age_limit', '41';
put 'threshold_ref', '8', 'limit:high_age_limit', '100';
put 'threshold_ref', '8', 'limit:low_value', '66';
put 'threshold_ref', '8', 'limit:high_value', '73';
put 'threshold_ref', '8', 'alert:alert_flag', '0';
put 'threshold_ref', '8', 'alert:alert_message', 'Normal';

put 'threshold_ref', '9', 'attribute:attribute', 'heartBeat';
put 'threshold_ref', '9', 'limit:low_age_limit', '41';
put 'threshold_ref', '9', 'limit:high_age_limit', '100';
put 'threshold_ref', '9', 'limit:low_value', '74';
put 'threshold_ref', '9', 'limit:high_value', '9999';
put 'threshold_ref', '9', 'alert:alert_flag', '1';
put 'threshold_ref', '9', 'alert:alert_message', 'Higher Heart Rate than Normal';

put 'threshold_ref', '10', 'attribute:attribute', 'bp';
put 'threshold_ref', '10', 'limit:low_age_limit', '41';
put 'threshold_ref', '10', 'limit:high_age_limit', '100';
put 'threshold_ref', '10', 'limit:low_value', '0';
put 'threshold_ref', '10', 'limit:high_value', '150';
put 'threshold_ref', '10', 'alert:alert_flag', '1';
put 'threshold_ref', '10', 'alert:alert_message', 'Low BP than Normal';

put 'threshold_ref', '11', 'attribute:attribute', 'bp';
put 'threshold_ref', '11', 'limit:low_age_limit', '41';
put 'threshold_ref', '11', 'limit:high_age_limit', '100';
put 'threshold_ref', '11', 'limit:low_value', '151';
put 'threshold_ref', '11', 'limit:high_value', '180';
put 'threshold_ref', '11', 'alert:alert_flag', '0';
put 'threshold_ref', '11', 'alert:alert_message', 'Normal';

put 'threshold_ref', '12', 'attribute:attribute', 'bp';
put 'threshold_ref', '12', 'limit:low_age_limit', '41';
put 'threshold_ref', '12', 'limit:high_age_limit', '100';
put 'threshold_ref', '12', 'limit:low_value', '181';
put 'threshold_ref', '12', 'limit:high_value', '9999';
put 'threshold_ref', '12', 'alert:alert_flag', '1';
put 'threshold_ref', '12', 'alert:alert_message', 'Higher BP than Normal';
```

Screenshot for insertion of records

```

hbase(main):003:0* put 'threshold_ref', '1', 'attribute:attribute', 'heartBeat';
put 'threshold_ref', '2', 'limit:high_value', '78';
put 'threshold_ref', '2', 'alert:alert_flag', '0';
put 'threshold_ref', '2', 'alert:alert_message', 'Normal';

put 'threshold_ref', '3', 'attribute:attribute', 'heartBeat';
put 'threshold_ref', '3', 'limit:low_age_limit', '0';
put 'threshold_ref', '3', 'limit:high_age_limit', '40';
put 'threshold_ref', '3', 'limit:low_value', '79';
put 'threshold_ref', '3', 'limit:high_value', '9999';
put 'threshold_ref', '3', 'alert:alert_flag', '1';
put 'threshold_ref', '3', 'alert:alert_message', 'Higher Heart Rate than Normal'
;

put 'threshold_ref', '4', 'attribute:attribute', 'bp';
put 'threshold_ref', '4', 'limit:low_age_limit', '0';
put 'threshold_ref', '4', 'limit:high_age_limit', '40';
put 'threshold_ref', '4', 'limit:low_value', '0';
put 'threshold_ref', '4', 'limit:high_value', '160';
put 'threshold_ref', '4', 'alert:alert_flag', '1';
put 'threshold_ref', '4', 'alert:alert_message', 'Low BP than Normal';

put 'threshold_ref', '5', 'attribute:attribute', 'bp';
put 'threshold_ref', '5', 'limit:low_age_limit', '0';
put 'threshold_ref', '5', 'limit:high_age_limit', '40';
put 'threshold_ref', '5', 'limit:low_value', '161';
put 'threshold_ref', '5', 'limit:high_value', '220';
put 'threshold_ref', '5', 'alert:alert_flag', '0';

```

View records in HBase table

```
scan 'threshold_ref';
```

```
hbase(main):098:0* scan 'threshold_ref'
```

Screenshots of records

ROW	COLUMN+CELL
1	column=alert:alert_flag, timestamp=1679727530966, value=1
1	column=alert:alert_message, timestamp=1679727530971, value=Low Heart Rate than Normal
1	column=attribute:attribute, timestamp=1679727530926, value=heartBeat
1	column=limit:high_age_limit, timestamp=1679727530951, value=40
1	column=limit:high_value, timestamp=1679727530962, value=69
1	column=limit:low_age_limit, timestamp=1679727530947, value=0
1	column=limit:low_value, timestamp=1679727530958, value=0
10	column=alert:alert_flag, timestamp=1679727531306, value=1
10	column=alert:alert_message, timestamp=1679727531310, value=Low BP than Normal
10	column=attribute:attribute, timestamp=1679727531291, value=bp
10	column=limit:high_age_limit, timestamp=1679727531297, value=100
10	column=limit:high_value, timestamp=1679727531303, value=150
10	column=limit:low_age_limit, timestamp=1679727531294, value=41
10	column=limit:low_value, timestamp=1679727531300, value=0
11	column=alert:alert_flag, timestamp=1679727531326, value=0
11	column=alert:alert_message, timestamp=1679727531329, value=Normal
11	column=attribute:attribute, timestamp=1679727531312, value=bp
11	column=limit:high_age_limit, timestamp=1679727531318, value=100
11	column=limit:high_value, timestamp=1679727531323, value=180
11	column=limit:low_age_limit, timestamp=1679727531315, value=41
11	column=limit:low_value, timestamp=1679727531321, value=151

```

8      column=alert:alert_message, timestamp=1679992681690, value
      =Normal
8      column=attribute:attribute, timestamp=1679992681657, value
      =heartBeat
8      column=limit:high_age_limit, timestamp=1679992681668, valu
      e=100
8      column=limit:high_value, timestamp=1679992681679, value=73
8      column=limit:low_age_limit, timestamp=1679992681663, value
      =41
8      column=limit:low_value, timestamp=1679992681674, value=66
9      column=alert:alert_flag, timestamp=1679992681728, value=1
9      column=alert:alert_message, timestamp=1679992681734, value
      =Higher Heart Rate than Normal
9      column=attribute:attribute, timestamp=1679992681696, value
      =heartBeat
9      column=limit:high_age_limit, timestamp=1679992681708, valu
      e=100
9      column=limit:high_value, timestamp=1679992681722, value=99
9      column=limit:low_age_limit, timestamp=1679992681702, value
      =41
9      column=limit:low_value, timestamp=1679992681715, value=74
12 row(s) in 0.3570 seconds

```

Threshold_Reference_Table in Hive

```

CREATE EXTERNAL TABLE Threshold_Reference_Table (
    key int,
    Attribute string,
    low_age_limit int,
    high_age_limit int,
    Low_Range_Value int,
    High_Range_Value int,
    Alert_Flag int,
    Alert_Message string
)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH SERDEPROPERTIES (
    'hbase.columns.mapping' = ':key, attribute:attribute, limit:low_age_limit, limit:high_age_limit, limit:low_value,
    limit:high_value, alert:alert_flag, alert:alert_message',
    'hbase.table.name' = 'threshold_ref'
)
TBLPROPERTIES ('hbase.mapred.output.outputtable' = 'threshold_ref');

```

Set up for the Hive and HBase integration

HBase Cluster

Cluster: Hbasecluster **Waiting** Cluster ready to run steps.

Summary

Application user interfaces

Monitoring

Hardware

Configurations

Events

Steps

Bootstrap actions

Summary

Configuration details

Summary

Configuration details

Application user interfaces

Network and hardware

Security and access

Summary

Configuration details

Application user interfaces

Network and hardware

Security and access

For the Hive-HBase integration on different clusters, few inbound rules were added to the security group for HBase master node and Hive master node.

Screenshot of HBase cluster's master node (Security Group – master rules)

Following rule is added:

Type: "Custom TCP Rule"

Protocol: "TCP"

Port Range: "16000 - 60000"

Source: "Custom" and enter the Master Security Group for Hive cluster's master node.

EC2 > Security Groups > sg-0a4b13fe01f78e60c - ElasticMapReduce-master > Edit inbound rules

Edit inbound rules

Inbound rules control the incoming traffic that's allowed to reach the instance.

Inbound rules

Security group rule ID	Type	Protocol	Port range	Source	Description - optional
sg-0b548d3004905ba08	Custom TCP	TCP	16000 - 6000	Custom	

Screenshot of Hive cluster's master node (Security Group – master rules)

Following rule is added:

Type: "Custom TCP Rule"

Protocol: "TCP"

Port: "10000"

Source: "Custom" and enter the Master Security Group for HBase cluster's master node.

sg-0f1b8be6ba73b57d7

Custom TCP

TCP

10000

Custom

Q

Delete

In the Hive Shell

- Connect the HBase client on your Hive cluster to the HBase cluster that contains your data.

set hbase.zookeeper.quorum= <public DNS name of the master node of the HBase cluster>;

```
set hbase.zookeeper.quorum=ec2-52-23-186-126.compute-1.amazonaws.com;
```

```
^C[hadoop@ip-172-31-82-68 ~]$ hive
Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.
properties Async: true
hive> set hbase.zookeeper.quorum=ec2-52-23-186-126.compute-1.amazonaws.com;
```

- Use database patient_health_care

```
use patient_health_care;
```

```
hive> use patient_health_care;
OK
Time taken: 0.046 seconds
```

- Create external table named *Threshold_Reference_Table*

```
CREATE EXTERNAL TABLE Threshold_Reference_Table (
  key int,
  Attribute string,
  low_age_limit int,
  high_age_limit int,
  Low_Range_Value int,
  High_Range_Value int,
  Alert_Flag int,
  Alert_Message string
)
STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
WITH SERDEPROPERTIES (
  'hbase.columns.mapping' = ':key, attribute:attribute, limit:low_age_limit, limit:high_age_limit,
limit:low_value, limit:high_value, alert:alert_flag, alert:alert_message',
  'hbase.table.name' = 'threshold_ref'
)
TBLPROPERTIES ('hbase.mapred.output.outputtable' = 'threshold_ref');
```

```
hive> CREATE EXTERNAL TABLE Threshold_Reference_Table (
  >   key int,
  >   Attribute string,
  >   low_age_limit int,
  >   high_age_limit int,
  >   Low_Range_Value int,
  >   High_Range_Value int,
  >   Alert_Flag int,
  >   Alert_Message string
  > )
  > STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'
  > WITH SERDEPROPERTIES (
  >   'hbase.columns.mapping' = ':key, attribute:attribute, limit:low_age_limit, limit:high_age_limit, limit:low_value, limit:high_value, alert:alert_flag, alert:alert_message',
  >   'hbase.table.name' = 'threshold_ref'
  > )
  > TBLPROPERTIES ('hbase.mapred.output.outputtable' = 'threshold_ref');
OK
Time taken: 2.31 seconds
```

- View the contents of Threshold_Reference_Table

```
set hive.cli.print.header = true;

SELECT * FROM Threshold_Reference_Table order by key;
```

```
hive> set hive.cli.print.header = true;
hive> select * from Threshold_Reference_Table order by key;
Query ID = hadoop_20230328084446_4e04390b-9e15-4799-a8f7-ef03baf510dd
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1679992808558_0001)

Map 1: -/-      Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 0/1      Reducer 2: 0/1
Map 1: 0(+1)/1  Reducer 2: 0/1
Map 1: 1/1      Reducer 2: 0(+1)/1
Map 1: 1/1      Reducer 2: 1/1
```

Screenshot of Threshold_Reference_Table records:

```
OK
threshold_reference_table.key    threshold_reference_table.attribute    threshold
d_reference_table.low_age_limit threshold_reference_table.high_age_limit t
hreshold_reference_table.low_range_value    threshold_reference_table.high_r
ange_value    threshold_reference_table.alert_flag    threshold_reference_tabl
e.alert_message
1      heartBeat      0      40      0      69      1      Low Heart Rate t
han Normal
2      heartBeat      0      40      70      78      0      Normal
3      heartBeat      0      40      79      9999      1      Higher Heart Rat
e than Normal
4      bp      0      40      0      160      1      Low BP than Normal
5      bp      0      40      161      220      0      Normal
6      bp      0      40      221      9999      1      Higer BP than Normal
7      heartBeat      41      100      0      65      1      Low Heart Rate t
han Normal
8      heartBeat      41      100      66      73      0      Normal
9      heartBeat      41      100      74      9999      1      Higher Heart Rat
e than Normal
10     bp      41      100      0      150      1      Low BP than Normal
11     bp      41      100      151      180      0      Normal
12     bp      41      100      181      9999      1      Higher BP than Normal
Time taken: 15.749 seconds, Fetched: 12 row(s)
```

- Create copy of threshold table *Threshold_Reference* in Hive and insert the records from Threshold_Reference_Table to Threshold_Reference

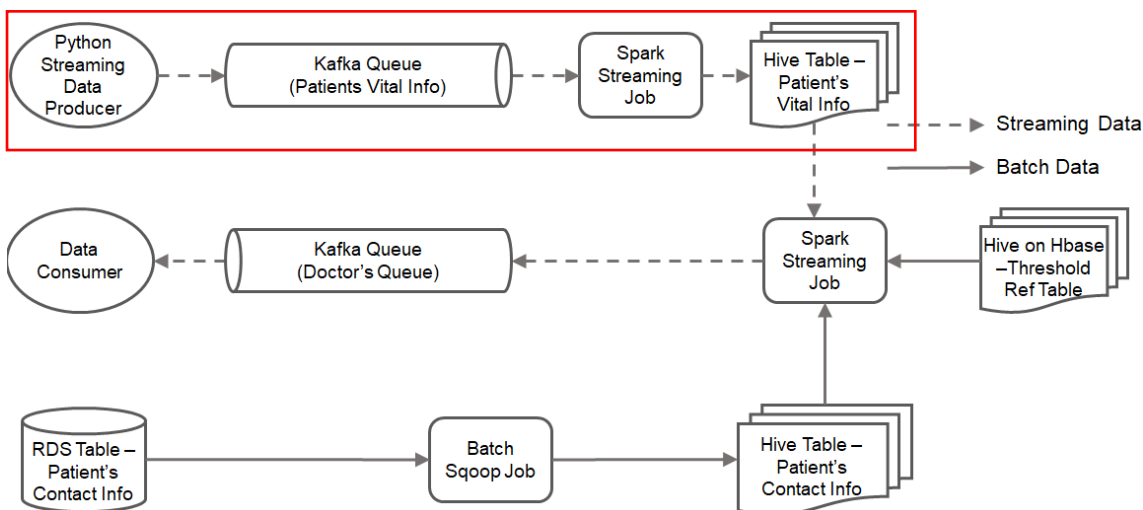
This table can be accessed by spark streaming application 2 for Part 4 mentioned below)

```
CREATE EXTERNAL TABLE Threshold_Reference (
  key int,
  Attribute string,
  low_age_limit int,
  high_age_limit int,
  Low_Range_Value int,
  High_Range_Value int,
  Alert_Flag int,
  Alert_Message string
);

INSERT INTO table Threshold_Reference SELECT * FROM Threshold_Reference_Table;
```

PART 3:

Taking streaming data (patients' vital information) and storing it in a table



This involves following tasks:

1. Stimulate streaming data by building a Kafka Producer application in Python to read data from RDS per second. This Kafka producer will push the patient vitals to the Kafka topic 'patient_vital_topic'
2. Write a spark streaming job to read data from Kakfa topic and add the timestamp column and store in HDFS location in parquet format
3. Create an external Hive table 'Patients_Vital_Info' that reads streaming data from HDFS location

There are 1800 records, therefore, the producer will be able to read all the records in 30 minutes.

Kafka EC2 Setup

1. Launch EC2 Instance
2. In the Application and OS Images section, you will need to select the Image to be used for the EC2 instance.
Click on "Browse more AMIs".

3. In the search box that appears at the top, copy and paste the following AMI id and press enter: ami-06c41d8b5a6ddd3c2

In the “Community AMIs” tab, you will find the AMI with the following AMI Name: Kafka_Anaconda-New-2022

Click on the “Select” button to choose the image.

4. Select the General Purpose m4.large type EC2 instance, as shown in the image below.

5. Select the Key Pair to login to the instance via SSH

6. In the “Network Settings” section, go with the default security group. Make sure that the option “Allow SSH traffic from Anywhere” is ticked. This will ensure that you’re able to SSH into the instance from your SSH client

7. In the “Configure storage” settings, you need to enter the volume size as 30 GiB and volume type as standard (magnetic)

8. Once the settings have been updated, click on the “Launch Instance” button to create the instance.

9. Security -> Security groups -> edit inbound rules

Click on Add Rule button and configure the security group as shown below. Enter the following values as shown in the image below

Type: Custom TCP

Port Range: 8888

Source: Anywhere-IPv4

10. Similarly, you need to add the following port numbers: 2181, 9092, 9000, 8080

11. Next click on the Save rules.

12. Go to Elastic IP -> Associate IP with currently running instance

13. Login to EC2 machine ,

14. `cd /home/ec2-user/downloads/kafka_2.12-2.3.0`

15. `cd config/`

`vi server.properties`

there is a line 36

`#advertised.listeners=PLAINTEXT://your.host.name:9092`

Then go to the above line and uncomment it by removing the #. Next in place of your.host.name enter the IPv4 Public IP of your EC2 instance. In this case, that line would read as follows:

`advertised.listeners=PLAINTEXT://52.21.15.133:9092`

`:wq!`

- Start Zookeeper server

```
cd downloads/kafka_2.12-2.3.0
bin/zookeeper-server-start.sh config/zookeeper.properties
```

- Start Kafka server

```
cd downloads/kafka_2.12-2.3.0
bin/kafka-server-start.sh config/server.properties
```

- Delete the topic if it already exists and create again

```
cd downloads/kafka_2.12-2.3.0

bin/kafka-topics.sh --bootstrap-server ec2-44-196-94-216.compute-1.amazonaws.com:9092 --delete --
topic patient_vital_topic
```

```
bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic patient_vital_topic
```

- Transfer the Python producer file and Install required packages

```
WINSXP transfer file to downloads/kafka_2.12-2.3.0 folder
```

```
pip install mysql-connector-python  
pip install mysql-connector-repackaged
```

- Run the producer script

```
python kafka_produce_patient_vitals.py
```

Spark Streaming Job 1 (*kafka_spark_patient_vitals.py*)

It reads data from Kafka topic and add the timestamp column and store in HDFS location in parquet format.

```
export SPARK_KAFKA_VERSION=0.10  
spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5 kafka_spark_patient_vitals.py
```

Note: In Spark streaming code, we have added a new column using `lit('2022-03-16')`, and while writing the data we have partitioned the data by this newly added constant column, thus all of the parquet files get written inside this folder `/user/livy/output/date=2022-03-16`.

This was done because `_spark_metadata` directory is by default created on the HDFS output directory. In our case, it will get created on path `/user/livy/output`, and due to partition by functionality, another folder is created named `'2022-03-16'` inside which the parquet files are present.

In the Hive Shell

- Create external table 'Patients_Vital_Info'

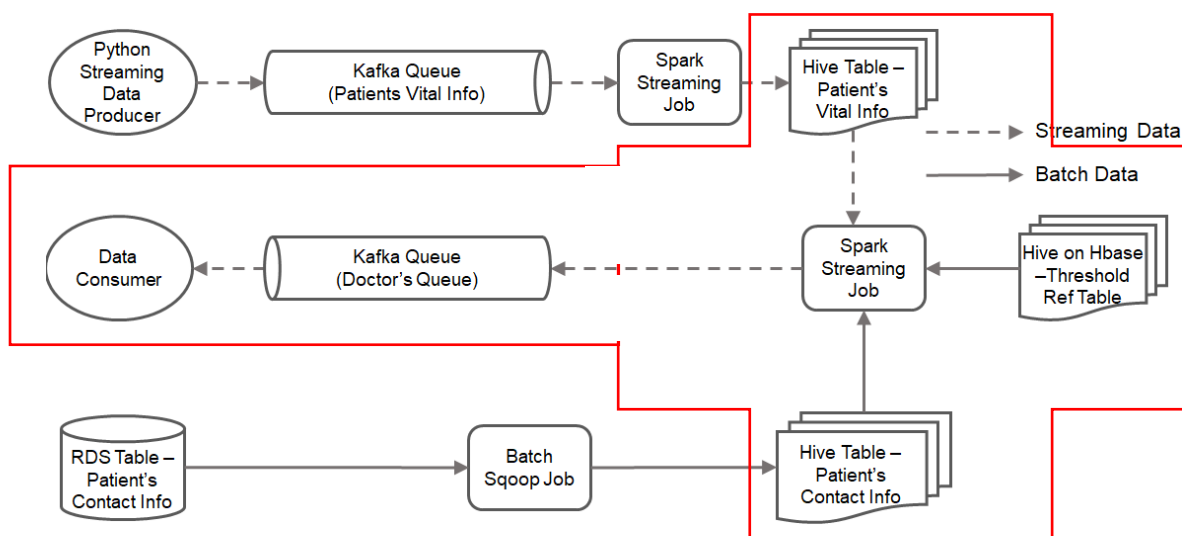
```
hive  
  
use patient_health_care;  
set hive.cli.print.header=true;  
  
CREATE EXTERNAL TABLE IF NOT EXISTS Patients_Vital_Info (  
    CustomerID int,  
    BP int,  
    HeartBeat int,  
    Message_time timestamp  
)  
STORED AS PARQUET  
LOCATION '/user/livy/output/date=2022-03-16';
```

```
[root@ip-172-31-82-23 ~]# hive

Logging initialized using configuration in file:/etc/hive/conf.dist/hive-log4j2.
properties Async: true
hive> create database if not exists patient_health_care;
OK
Time taken: 1.098 seconds
hive> Show databases;
OK
default
patient_health_care
Time taken: 0.323 seconds, Fetched: 2 row(s)
hive>
  > use patient_health_care;
OK
Time taken: 0.111 seconds
hive>
  >
  > CREATE EXTERNAL TABLE IF NOT EXISTS Patients_Vital_Info (
  >   CustomerID int,
  >   BP int,
  >   HeartBeat int,
  >   Message_time timestamp
  > )
  > STORED AS PARQUET
  > LOCATION '/user/livy/output/date=2022-03-16';
OK
Time taken: 0.333 seconds
```

PART 4:

Comparing the vital information with threshold information and analysing and sending notifications if the data is out of the threshold limits



This involves following tasks:

1. Create a Kafka topic named **Alerts_Message** to store irregular patient vitals.
2. Write spark streaming job to read data from 3 hive tables and analyse the patient vitals. If they are irregular, push them to a Kafka topic.
3. Create a Kafka consumer to read messages pushed to the above topic
4. Send email notifications for the messages read by consumer using SNS.

Create a Kafka topic (Alerts_Message)

- Delete the topic if it already exists

```
bin/kafka-topics.sh --bootstrap-server ec2-44-196-94-216.compute-1.amazonaws.com:9092 --delete --topic Alerts_Message
```

- Create the topic again

```
bin/kafka-topics.sh --create --bootstrap-server localhost:9092 --replication-factor 1 --partitions 1 --topic Alerts_Message
```

- View the list of topics

```
bin/kafka-topics.sh --list --bootstrap-server localhost:9092
```

Spark Streaming Job to push irregular patient vitals to 'Alerts_Message' Kafka topic

```
export SPARK_KAFKA_VERSION=0.10
spark-submit --packages org.apache.spark:spark-sql-kafka-0-10_2.11:2.4.5
kafka_spark_generate_alerts.py
```

Configure SNS

Create SNS topic (Health_Alerts)

Amazon SNS > Topics > Health_Alerts

Health_Alerts

Edit Delete Publish message

Details

Name Health_Alerts	Display name -
ARN arn:aws:sns:us-east-1:545120555452:Health_Alerts	Topic owner 545120555452
Type Standard	

[Subscriptions](#) | [Access policy](#) | [Data protection policy](#) | [Delivery policy \(HTTP/S\)](#) | [Delivery status logging](#) | [Encryption](#) | [Tags](#) | [Integrations](#)

Subscriptions (1) Edit Delete Request confirmation Confirm subscription **Create subscription**

ID	Endpoint	Status	Protocol
c08784f0-8d2f-4439-a5bc-7b9f3b671cc9	mehak.aggarwal135@gmail.com	Confirmed	EMAIL

Subscribe to SNS topic

Subscription: c08784f0-8d2f-4439-a5bc-7b9f3b671cc9

Edit

Delete

Details

ARN

arn:aws:sns:us-east-1:545120555452:Health_Alerts:c08784f0-8d2f-4439-a5bc-7b9f3b671cc9

Endpoint

mehak.aggarwal135@gmail.com

Topic

[Health_Alerts](#)

Subscription Principal

arn:aws:iam::545120555452:role/voclabs

Status

✔ Confirmed

Protocol

EMAIL

Subscription Confirmation Email

AWS Notification - Subscription Confirmation Inbox x**AWS Notifications** <no-reply@sns.amazonaws.com>
to me ▾

Sat, Apr 8, 10:53 PM (17 hours ago)



You have chosen to subscribe to the topic:

arn:aws:sns:us-east-1:545120555452:Health_Alerts

To confirm this subscription, click or visit the link below (If this was in error no action is necessary):

[Confirm subscription](#)Please do not reply directly to this email. If you wish to remove yourself from receiving all future SNS subscription confirmation requests please send an email to [sns-opt-out](#)

Reply

Forward

Kafka consumer

```
pip install kafka-python
pip install boto3
python kafka_consume_alerts.py
```

Final Output Screenshot:

Patient health Alert email - using SNS

Health Alert Notification Inbox x**AWS Notifications** <no-reply@sns.amazonaws.com>
to me ▾

1:50 AM (14 hours ago)



```
{
  "patientname": "Karan C",
  "age": 56,
  "patientaddress": "Aws Address",
  "phone_number": "8923739282",
  "admitted_ward": 3,
  "bp": 171,
  "heartBeat": 56,
  "input_message_time": "2023-04-08T17:58:12.698Z",
  "alert_message": "Low Heart Rate than Normal",
  "alert_generated_time": "2023-04-08T17:58:16.266Z"
}
```

--

If you wish to stop receiving notifications from this topic, please click or visit the link below to unsubscribe:

https://sns.us-east-1.amazonaws.com/unsubscribe.html?SubscriptionArn=arn:aws:sns:us-east-1:545120555452:Health_Alerts:c08784f0-8d2f-4439-a5bc-7b9f3b671cc9&Endpoint=mehak.aggarwal135@gmail.comPlease do not reply directly to this email. If you have any questions or comments regarding this email, please contact us at <https://aws.amazon.com/support>

Reply

Forward

In short

- Create HBase table, install sqoop and run import
- Then create all external hove tables
- Set up SNS configurations
- Run Kafka zookeeper, server, create topic, producer & consumer in separate SSH terminal windows
- Run both spark streaming jobs